

DIS Sheet 4

a) Describe how the program works and how the persistence manager accomplishes logging and recovery

- We have a persistence manager, who inhabits the three important functions "write", "commit" and "safe_persistent". Whenever a Client begins a transaction he gets his own transactionnumber (taid). Then when he gives a write in order this write gets his own lognumber (lsn) and is being safed with lsn, taid, pageid and data in the hashtable, which is our buffer, and the logdata. If the client wants to commit, this commit is safed with "EOT" in the logfile, but not in the buffer. Now it doesn't matter which client has given any writes in order, but as long as the amount of writes safed in the buffer exceeds 5, when a new write is being ordered, the buffer automatically checks whether the write orders in the buffer have already been committed. This is done by looking in the logfile and checking if there are any commits (can be seen by "EOT") for a specific taid. If there are commits for write commands in the buffer then we start the function safe_persistent. This function safes the write commands in the userdata. After a write command is safed persistent in the userdata we delete this writecommand in the buffer. If we don't find any commits, but the number of writes in the buffer is bigger than 5, it doesn't matter, since we will check again when a new write command is being ordered. At the end of the system all write operations with a commit are being safed, since we let the recoverymode also run at the end of the system.
- the recovery_mode starts whenever we start the program. It checks the logfile for write orders which haven't been safed persistent and also haven't been overwritten by any other write orders in the userdata (this is checked by comparing lsn in userdata and in logfile). For these write commands the system will then place a new write order with the same taid, pageid and data (only lsn will be new). By doing this the write commands will then be safed when a commit for the taid is/was placed. When we execute a write function from the recoverymode we don't have to wait till the buffer exceeds a number of entries, since we let the writefunction know we come from the recoverymode by giving him a Boolean.

- b) Tested logging and buffering (with screenshots)
i) uncommitted transactions in the log file

```
RecoveryTool.recoverymode();
Integer client_1_taid = client_1.beginTransaction();
Integer client_2_taid = client_2.beginTransaction();
Integer client_3_taid = client_3.beginTransaction();
Integer client_4_taid = client_4.beginTransaction();
Integer client_5_taid = client_5.beginTransaction();

client_1.write(client_1_taid, 15, "test1");
client_1.write(client_1_taid, 17, "test2");
client_2.write(client_2_taid, 24, "test3");
client_1.commit(client_1_taid);
client_2.write(client_2_taid, 28, "test4");
client_2.write(client_2_taid, 27, "test5");
client_2.write(client_2_taid, 26, "test6");
client_2.write(client_2_taid, 23, "test7");
// client_2.write(client_2_taid, 23, "test7 client_2 üb
```

```
≡ userdata.txt ● ≡ logdata.txt ×
Loading file tree... > ≡ logdata.txt
1 | 1, 1, 15, test1
2 | 2, 1, 17, test2
3 | 3, 2, 24, test3
4 | 4, 1, , EOT
5 | 5, 2, 28, test4
6 | 6, 2, 27, test5
7 | 7, 2, 26, test6
8 | 8, 2, 23, test7
9 |
```

```
≡ userdata.txt ● ≡ logdata.txt ×
Loading file tree... > ≡ userdata.txt
1 | 10, -1
2 | 11, -1
3 | 12, -1
4 | 13, -1
5 | 14, -1
6 | 15, 1, test1
7 | 16, -1
8 | 17, 2, test2
9 | 18, -1
10 | 19, -1
11 | 20, -1
12 | 21, -1
13 | 22, -1
14 | 23, -1
15 | 24, -1
16 | 25, -1
17 | 26, -1
18 | 27, -1
19 | 28, -1
20 | 29, -1
21 | 30, -1
22 | 31, -1
23 | 32, -1
24 | 33, -1
```

- c) Recovery
- i) Total program run (with changed pages)

≡ userdata.txt ●

≡ logdata.txt

⌵ Pers

Loading file tree... > ≡ userdata.txt

1

10, -1

2

11, -1

3

12, -1

4

13, -1

5

14, -1

6

15, 1, test1

7

16, -1

8

17, 2, test2

9

18, -1

10

19, -1

11

20, -1

12

21, -1

13

22, -1

14

23, 9, test7 client_2 überschreib

15

24, 3, test3

16

25, -1

17

26, 7, test6

18

27, 6, test5

19

28, 5, test4

14, -1

15, 1

16, -1

17, 48, test2 client_4 überschreibt client_1

18, -1

19, -1

20, -1

21, -1

22, -1

23, 39, test7 client_2 überschreibt sich selbst

24, 3, ich wurde geändert

14, -1

15, 31, test1

16, -1

17, 48, test2 client_4 überschreibt client_1

18, -1

19, -1

20, -1

21, -1

22, -1

23, 39, test7 client_2 überschreibt sich selbst

24, 33, test3

- ii) changed buffer size (current state of the page)
user data is still empty

and log data:

```
userdata.txt  logdata.txt  PersistenceManager.java
Loading file tree... > logdata.txt
1 1, 1, 15, test1
2 2, 1, 17, test2
3 3, 2, 24, test3
4 4, 1, , EOT
5 5, 2, 28, test4
6 6, 2, 27, test5
7 7, 2, 26, test6
8 8, 2, 23, test7
9 9, 2, 23, test7 client_2 überschreibt sich selbst
10 10, 2, , EOT
11 11, 3, 34, test8
12 12, 3, 36, test9
13 13, 3, 39, test10
14 14, 4, 45, test11
15 15, 3, , EOT
16 16, 4, 46, test13
17 17, 4, 47, test14
18 18, 4, 17, test2 client_4 überschreibt client_1
19 19, 4, , EOT
20 20, 5, 52, test15
21 21, 5, 54, test17
22 22, 5, 54, test18
23 23, 5, 54, test19
24 24, 5, 54, test20
25 25, 5, , EOT
26
```

Result:

```
userdata.txt  logdata.txt  PersistenceManager.java
Loading file tree... > logdata.txt
1 1, 1, 15, test1
2 2, 1, 17, test2
3 3, 2, 24, test3
4 4, 1, , EOT
5 5, 2, 28, test4
6 6, 2, 27, test5
7 7, 2, 26, test6
8 8, 2, 23, test7
9 9, 2, 23, test7 client_2 überschreibt sich selbst
10 10, 2, , EOT
11 11, 3, 34, test8
12 12, 3, 36, test9
13 13, 3, 39, test10
14 14, 4, 45, test11
15 15, 3, , EOT
16 16, 4, 46, test13
17 17, 4, 47, test14
18 18, 4, 17, test2 client_4 überschreibt client_1
19 19, 4, , EOT
20 20, 5, 52, test15
21 21, 5, 54, test17
22 22, 5, 54, test18
23 23, 5, 54, test19
24 24, 5, 54, test20
25 25, 5, , EOT
26 26, 1, 15, test1
27 27, 1, 17, test2
28 28, 4, 17, test2 client_4 überschreibt client_1
29 29, 2, 23, test7
30 30, 2, 23, test7 client_2 überschreibt sich selbst
31 31, 2, 24, test3

userdata.txt  logdata.txt  PersistenceManager.java
Loading file tree... > userdata.txt
1 10, -1
2 11, -1
3 12, -1
4 13, -1
5 14, -1
6 15, 28, test1
7 16, -1
8 17, 28, test2 client_4 überschreibt client_1
9 18, -1
10 19, -1
11 20, -1
12 21, -1
13 22, -1
14 23, 30, test7 client_2 überschreibt sich selbst
15 24, 31, test3
16 25, -1
17 26, 32, test6
18 27, 33, test5
19 28, 34, test4
20 29, -1
21 30, -1
22 31, -1
23 32, -1
24 33, -1
25 34, 35, test8
26 35, -1
27 36, 36, test9
28 37, -1
29 38, -1
30 39, 37, test10
31 40, -1
32 41, -1
33 42, -1
34 43, -1
35 44, -1
36 45, 38, test11
37 46, 39, test12
```

iii) changed buffer size (changed lsn)

```
≡ userdata.txt ● ≡ logdata.txt
Loading file tree... > ≡ userdata.txt
1      10, -1
2      11, -1
3      12, -1
4      13, -1
5      14, -1
6      15, -1
7      16, -1
8      17, -1
9      18, -1
10     19, -1
11     20, -1
12     21, -1
13     22, -1
14     23, 100
```

```
≡ userdata.txt ≡ logdata.txt × PersistenceManager.java
Loading file tree... > ≡ logdata.txt
1 + 1, 1, 15, test1
2   2, 1, 17, test2
3   3, 2, 24, test3
4   4, 1,   , EOT
5   5, 2, 28, test4
6   6, 2, 27, test5
7   7, 2, 26, test6
8   8, 2, 23, test7
9   9, 2, 23, test7 client_2 überschreibt sich selbst
10  10, 2,   , EOT
11  11, 3, 34, test8
12
```

Result: The user data is still empty at pageid 23

```
userdata.txt • logdata.txt
loading file tree... > userdata.txt
1      10, -1
2      11, -1
3      12, -1
4      13, -1
5      14, -1
6      15, 12, test1
7      16, -1
8      17, 13, test2
9      18, -1
10     19, -1
11     20, -1
12     21, -1
13     22, -1
14     23, 100
15     24, 14, test3
16     25, -1
17     26, 15, test6
18     27, 16, test5
19     28, 17, test4
20     29, -1
```

userdata for pageid 23 is empty because the write command didnt get pushed into the hashtable again, since the lsn in the page is higher than in the logdata